

Lyft-Uber-Price-Prediction

7 November 2019

IMPORTING DATASETS AND CLEANING THEM

Importing dataset cab_rides

```
cab_rides <- read.csv("C:/Users/nisht/Desktop/MITA/Fall/MVA/Final Project/cab_rides.csv")
summary(cab_rides)
```

```
##      distance      cab_type      time_stamp
##  Min.   :0.020    Lyft:307408    Min.     :1.543e+12
##  1st Qu.:1.280    Uber:385663    1st Qu.  :1.543e+12
##  Median :2.160                      Median   :1.544e+12
##  Mean   :2.189                      Mean     :1.544e+12
##  3rd Qu.:2.920                      3rd Qu.  :1.545e+12
##  Max.   :7.860                      Max.     :1.545e+12
##
##              destination              source              price
##  Financial District: 58851    Financial District: 58857    Min.    : 2.50
##  Theatre District  : 57798    Theatre District  : 57813    1st Qu.: 9.00
##  Back Bay          : 57780    Back Bay         : 57792    Median :13.50
##  Boston University : 57764    Boston University : 57764    Mean   :16.55
##  Haymarket Square  : 57764    North End        : 57763    3rd Qu.:22.50
##  Fenway            : 57757    Fenway           : 57757    Max.   :97.50
##  (Other)           :345357    (Other)          :345325    NA's   :55095
##  surge_multiplier      id
##  Min.   :1.000    00005b8c-5647-4104-9ac6-94fa6a40f3c3: 1
##  1st Qu.:1.000    00006eeb-0183-40c1-8198-c441d3c8a734: 1
##  Median :1.000    00008b42-5ecc-4f66-b4b9-b22a331634e6: 1
##  Mean   :1.014    000094c0-00c4-43f1-ae1b-4693eec2a580: 1
##  3rd Qu.:1.000    0000a8b2-e4d3-4227-8374-af8a2366e475: 1
##  Max.   :3.000    0000b5d6-59be-4534-b371-8214334d94f0: 1
##              (Other)              :693065
##              product_id              name
##  6d318bcc-22a3-4af6-bddd-b409bfce1546: 55096    Black SUV: 55096
##  6f72dfc5-27f1-42e8-84db-ccc7a75f6969: 55096    UberXL   : 55096
##  9a0e7b09-b92b-4c41-9779-2ad22b4d779d: 55096    WAV      : 55096
##  6c84fd89-3f11-4782-9b50-97c468b19529: 55095    Black    : 55095
##  8cf7e821-f0d3-49c6-8eba-e679c0ebcf6a: 55095    Taxi     : 55095
##  55c66225-fbe7-4fd5-9072-eab1ece5e23e: 55094    UberX    : 55094
##  (Other)              :362499    (Other)  :362499
```

```
cab_data<-cab_rides
```

Creating a date_time column

```
cab_data$date_time<-as.POSIXct((cab_data$time_stamp/1000),origin = "1970-01-01 00:53:20", tz="GMT")
```

Importing dataset weather

```
weather <- read.csv("C:/Users/nisht/Desktop/MITA/Fall/MVA/Final Project/weath  
er.xls")
```

```
summary(weather)
```

```
##      i..temp      location      clouds
##  Min.   :19.62    Back Bay      : 523    Min.   :0.0000
##  1st Qu.:36.08    Beacon Hill    : 523    1st Qu.:0.4400
##  Median :40.13    Boston University : 523    Median :0.7800
##  Mean   :39.09    Fenway         : 523    Mean   :0.6778
##  3rd Qu.:42.83    Financial District: 523    3rd Qu.:0.9700
##  Max.   :55.41    Haymarket Square : 523    Max.   :1.0000
##              (Other)      :3138
##      pressure      rain      time_stamp      humidity
##  Min.   : 988.2    Min.   :0.000    Min.   :1.543e+09    Min.   :0.450
##  1st Qu.: 997.7    1st Qu.:0.005    1st Qu.:1.543e+09    1st Qu.:0.670
##  Median :1007.7    Median :0.015    Median :1.544e+09    Median :0.760
##  Mean   :1008.4    Mean   :0.058    Mean   :1.544e+09    Mean   :0.764
##  3rd Qu.:1018.5    3rd Qu.:0.061    3rd Qu.:1.545e+09    3rd Qu.:0.890
##  Max.   :1035.1    Max.   :0.781    Max.   :1.545e+09    Max.   :0.990
##              NA's      :5382
##      wind
##  Min.   : 0.290
##  1st Qu.: 3.518
##  Median : 6.570
##  Mean   : 6.803
##  3rd Qu.: 9.920
##  Max.   :18.180
##
```

```
str(weather)
```

```
## 'data.frame':    6276 obs. of  8 variables:
##  $ i..temp      : num  42.4 42.4 42.5 42.1 43.1 ...
##  $ location     : Factor w/ 12 levels "Back Bay","Beacon Hill",...: 1 2 3 4 5
##  $ clouds       : num  1 1 1 1 1 1 1 1 1 1 ...
##  $ pressure     : num  1012 1012 1012 1012 1012 ...
##  $ rain         : num  0.1228 0.1846 0.1089 0.0969 0.1786 ...
##  $ time_stamp: int  1545003901 1545003901 1545003901 1545003901 1545003901
##  $ humidity     : num  0.77 0.76 0.76 0.77 0.75 0.77 0.77 0.77 0.78 0.75 ...
##  $ wind         : num  11.2 11.3 11.1 11.1 11.5 ...
```

```
weather_data<-weather
```

creating a date_time column in weather_data

```
weather_data$date_time<-as.POSIXct(weather_data$time_stamp,origin = "1970-01-01 00:53:20", tz="GMT")
str(weather_data)

## 'data.frame':    6276 obs. of  9 variables:
## $ i..temp      : num  42.4 42.4 42.5 42.1 43.1 ...
## $ location     : Factor w/ 12 levels "Back Bay","Beacon Hill",...: 1 2 3 4 5 6 7 8 9 10 ...
## $ clouds       : num   1 1 1 1 1 1 1 1 1 1 ...
## $ pressure     : num  1012 1012 1012 1012 1012 ...
## $ rain         : num   0.1228 0.1846 0.1089 0.0969 0.1786 ...
## $ time_stamp: int   1545003901 1545003901 1545003901 1545003901 1545003901 1545003901 1545003901 1545003901 1545003901 ...
## $ humidity    : num   0.77 0.76 0.76 0.77 0.75 0.77 0.77 0.77 0.78 0.75 ...
## $ wind        : num   11.2 11.3 11.1 11.1 11.5 ...
## $ date_time   : POSIXct, format: "2018-12-17 00:38:21" "2018-12-17 00:38:21" ...
```

merge the datasets to reflect the same time for a location

```
cab_data$merge_date<-paste(cab_data$source,"-",as.Date(cab_data$date_time),"-",format(cab_data$date_time,"%H:%M:%S"))
weather_data$merge_date<-paste(weather_data$location,"-",as.Date(weather_data$date_time),"-",format(weather_data$date_time,"%H:%M:%S"))

#making those values as characters
weather_data$merge_date<-as.character(weather_data$merge_date)
cab_data$merge_date<-as.character(cab_data$merge_date)
```

verify that merge_date has unique values.

```
weather_data<-subset(weather_data,!duplicated(weather_data$merge_date))
isTRUE(duplicated(weather_data$merge_date))

## [1] FALSE
```

Merging both the dataframes.

```
merge_data<-merge(x=weather_data, y=cab_data,by='merge_date', all.x=TRUE)
#str(merge_data)

merge_data$rain<-as.numeric(merge_data$rain)
merge_data$rain[is.na(merge_data$rain)]<-0

for ( i in 1:length(merge_data$rain)){
  if(merge_data$rain[i]>0 & merge_data$rain[i]<=0.30){
    merge_data$rain[i]=1
  }
}
```

```

}

for ( i in 1:length(merge_data$rain)){
  if(merge_data$rain[i]>=0.30 & merge_data$rain[i]!=1){
    merge_data$rain[i]=2
  }
}

merge_data$rain = factor(merge_data$rain,
                        levels = c(0,1,2),
                        labels = c(0,1,2))

```

Handling Missing values

```

#Extracting the numerical columns in a new dataframe "df"
merge_data$temp<-merge_data[,c(2)] #renaming a column
df<-merge_data[,c(4,5,8,9,10,11,22,16)]

#Data preparation
#Dealing with missing values
#summary(merge_data)
#summary(df)

merge_data$distance = ifelse(is.na(merge_data$distance),
                             ave(merge_data$distance , FUN = function
(x) mean(x, na.rm = TRUE))),
                             merge_data$surge_multiplier)

merge_data$price = ifelse(is.na(merge_data$price),
                           ave(merge_data$price , FUN = function(x) mean(x, na
.rm = TRUE))),
                           merge_data$price)

df$distance = ifelse(is.na(df$distance),
                     ave(df$distance , FUN = function(x) mean(x, na.rm = TRUE
))),
                     df$distance)

df$price = ifelse(is.na(df$price),
                  ave(df$price , FUN = function(x) mean(x, na.rm = TRUE))),
                  df$price)

```

Checking for null values

```
any(is.na(df))
```

```
## [1] FALSE
```

Adding date and time column in the df data set

```
df$day<-weekdays(df$date_time)
df$time<-format(df$date_time.x,"%H:%M:%S")
df$date_time<-as.Date(df$date_time.x)
merge_data$day=weekdays(merge_data$date_time.x)
```

Creating a Numeric dataframe

```
x<-df[,c(1,2,3,4,6,7,8)]
head(x)

##   clouds pressure humidity wind distance temp price
## 1    0.87  1014.39     0.92  1.46  2.168125 41.04 16.67376
## 2    0.86  1014.17     0.93  2.57  2.168125 40.63 16.67376
## 3    0.86  1014.17     0.93  2.59  1.440000 40.63  8.50000
## 4    0.86  1014.17     0.93  2.65  1.360000 40.61 16.50000
## 5    0.86  1014.17     0.93  2.65  1.220000 40.61 16.67376
## 6    0.95  1013.78     0.92  2.59  1.340000 40.72 26.50000
```


Let's check for multivariate analysis using chi-square plot

CORRELATION, COVARIANCE AND DISTANCE

```
#We are calculating for: clouds, pressure, rain, humidity, wind, distance, temp
covariance<-cov(x) #variance-covariance matrix created
correlation<-cor(x) #standardized
#colmeans
cm<-colMeans(x)
distance<-dist(scale(x,center=FALSE))
#Calculating di(generalized distance for all observations of our data)
d <- apply(x, MARGIN = 1, function(x) + t(x - cm) %*% solve(covariance) %*% (
x - cm))
```

The sorted distance are now plotted against the appropriate quantiles of the chi-distribution

```
plot(qc <- qchisq((1:nrow(x) - 1/2) / nrow(x), df = 5), sd <- sort(d),xlab =
expression(paste(chi[5]^2, " Quantile")),ylab = "Ordered distances")
oups <- which(rank(abs(qc - sd), ties = "random") > nrow(x) - 5)
text(qc[oups], sd[oups] - 1.5,oups)
abline(a=0,b=1,col="red")
```




#Our observations seems to deviate from linearity after a certain point

There is a complete deviation from Normality. We will apply the log transformation on our dataset.

```
#x_new<-x+1
#x_new=log(x - (min(x) - 1))
x_new<-log(x+1)

covariance<-cov(x_new) #variance-covariance matrix created
correlation<-cor(x_new) #standardized
#colmeans
cm<-colMeans(x_new)
distance<-dist(scale(x_new,center=FALSE))
#Calculating di(generalized distance for all observations of our data)
d <- apply(x_new, MARGIN = 1, function(x_new) + t(x_new - cm) %*% solve(covariance) %*% (x_new - cm))

plot(qc <- qchisq((1:nrow(x_new) - 1/2) / nrow(x_new), df = 6), sd <- sort(d),
     xlab = expression(paste(chi[6]^2, " Quantile")),ylab = "Ordered distances")
oups <- which(rank(abs(qc - sd), ties = "random") > nrow(x) - 6)
text(qc[oups], sd[oups] - 1.5,oups)
abline(a=0,b=1,col="red")
```



We have normalized the data..

Pca || T-test || F-test

Get the Correlations between the measurements

```
x_new<-x_new[,-7]
cor(x_new)

##           clouds      pressure      humidity      wind      distance
## clouds      1.00000000  0.04290237  0.380934772 -0.03065604  0.018863302
## pressure    0.04290237  1.00000000  0.051631425 -0.57889122  0.063942888
## humidity    0.38093477  0.05163143  1.000000000 -0.36402079  0.007197305
## wind        -0.03065604 -0.57889122 -0.364020786  1.00000000 -0.026571018
## distance    0.01886330  0.06394289  0.007197305 -0.02657102  1.000000000
## temp        0.51904189 -0.18953119  0.340366261  0.11773558 -0.003364190
##
##           temp
## clouds      0.51904189
## pressure    -0.18953119
## humidity    0.34036626
## wind        0.11773558
## distance    -0.00336419
## temp        1.00000000

sapply(x_new, sd, na.rm = TRUE)
```

```
## clouds pressure humidity wind distance temp
## 0.1997804 0.0124154 0.0686079 0.5304428 0.2347243 0.1437917
```

*#There are not considerable differences between these standard deviations.. S
till Let's see the PCAs.*


Let's Visualize Correlation..

```
library(corrplot)
```

```
## Warning: package 'corrplot' was built under R version 3.5.3
```

```
## corrplot 0.84 loaded
```

```
corrplot(cor(x_new), method="ellipse")
```



Using prcomp to compute the principal components (eigenvalues and eigenvectors).

With scale=TRUE, variable means are set to zero, and variances set to one

```
x_pca <- prcomp(x_new, scale=TRUE)
#x_pca$rotation
summary(x_pca)
```

```
## Importance of components:
##
##          PC1      PC2      PC3      PC4      PC5      PC6
## Standard deviation  1.3666 1.3004 0.9992 0.8339 0.6668 0.55060
## Proportion of Variance 0.3113 0.2818 0.1664 0.1159 0.0741 0.05053
## Cumulative Proportion 0.3113 0.5931 0.7595 0.8754 0.9495 1.00000
```

```
#x_pca$rotation
```

Each of these explains a percent of total variation in the dataset. PC1 explains 27.8% of total variance, PC2 explains 26% of total variance.. we need to go to PC5 to get a very accurate view on where it stands in relation to other samples as PC1-PC5 can explain 89.9% of the variance. sample scores stored in `x_pca$singularvalues` (square roots of eigenvalues) stored in `x_pca$dev` loadings (eigenvectors) are stored in `x_pca$rotation` variable means stored in `x_pca$center` standard deviations stored in `x_pca$scale` A table containing eigenvalues and %'s accounted, follows Eigenvalues are $sdev^2$

```
eigen_x <- x_pca$sdev^2
names(eigen_x) <- paste("PC", 1:6, sep="")
#eigen_x
sumlambdas <- sum(eigen_x)
sumlambdas #total sample variance

## [1] 6

propvar <- eigen_x/sumlambdas
#propvar
cumvar_x <- cumsum(propvar)
#cumvar_x
matlambdas <- rbind(eigen_x, propvar, cumvar_x)
rownames(matlambdas) <- c("Eigenvalues", "Prop. variance", "Cum. prop. variance")
round(matlambdas, 4)

##          PC1      PC2      PC3      PC4      PC5      PC6
## Eigenvalues  1.8675 1.6910 0.9983 0.6954 0.4446 0.3032
## Prop. variance 0.3113 0.2818 0.1664 0.1159 0.0741 0.0505
## Cum. prop. variance 0.3113 0.5931 0.7595 0.8754 0.9495 1.0000
```

Sample scores stored in x_pca\$x

We need to calculate the scores on each of these components for each individual in our sample.

```
#x_pca$rotation
xtyp_pca <- cbind(data.frame(df$price),x_pca$x)
str(xtyp_pca)

## 'data.frame':  9306 obs. of  7 variables:
## $ df.price: num  16.7 16.7 8.5 16.5 16.7 ...
## $ PC1      : num  -2.03 -1.81 -1.76 -1.75 -1.74 ...
## $ PC2      : num  -1.501 -1.095 -1.001 -0.972 -0.952 ...
## $ PC3      : num   0.1432 0.0861 1.187 1.3248 1.5828 ...
## $ PC4      : num  -0.474 -0.428 -0.309 -0.288 -0.261 ...
## $ PC5      : num   0.1565 -0.0283 -0.0348 -0.0438 -0.0449 ...
## $ PC6      : num  -0.56 -0.05892 -0.02739 -0.00333 0.0024 ...

#xtyp_pca
```


Merging price column

```
colnames(xtyp_pca)[colnames(xtyp_pca)=="df.price"] <- "price"
str(xtyp_pca)

## 'data.frame':  9306 obs. of  7 variables:
## $ price: num  16.7 16.7 8.5 16.5 16.7 ...
## $ PC1  : num  -2.03 -1.81 -1.76 -1.75 -1.74 ...
## $ PC2  : num  -1.501 -1.095 -1.001 -0.972 -0.952 ...
## $ PC3  : num   0.1432 0.0861 1.187 1.3248 1.5828 ...
## $ PC4  : num  -0.474 -0.428 -0.309 -0.288 -0.261 ...
## $ PC5  : num   0.1565 -0.0283 -0.0348 -0.0438 -0.0449 ...
## $ PC6  : num  -0.56 -0.05892 -0.02739 -0.00333 0.0024 ...
```

Plotting the scores of Pricipal Component 1 and Principal component 2


```
plot(xtyp_pca$PC1, xtyp_pca$PC2,xlab="PC1:", ylab="PC2")
abline(h=0)
abline(v=0)
```



Plotting the Variance of Principal Components

```
plot(eigen_x, xlab = "Component number", ylab = "Component variance", type =  
"b", main = "Scree diagram")
```

Scree diagram




#Plotting the Log

variance of COmponents

```
plot(log(eigen_x), xlab = "Component number", ylab = "log(Component variance)"  
, type="o", main = "Log(eigenvalue) diagram")
```


Log(eigenvalue) diagram



#Cumulative scree

plot

```
plot(cumsum(propvar), xlab = "Principal Component",  
     ylab = "Cumulative Proportion of Variance Explained",  
     type = "b")
```



Variance of the principal components


```
#View(x_pca)
diag(cov(x_pca$x))

##      PC1      PC2      PC3      PC4      PC5      PC6
## 1.8675135 1.6909966 0.9983146 0.6953850 0.4446266 0.3031637

#x_pca$x[,1]
#x_pca$x
```

Plotting the scores


```
xlim <- range(x_pca$x[,1])
plot(x_pca$x,xlim=xlim,ylim=xlim)
```



```
#x_pca$rotation[,1]  
#x_pca$rotation
```

Variance plot for each component. We can see that all components play a dominant role.

```
plot(x_pca)
```



#Taking first 4

components

```
xtyp_pca<-xtyp_pca[1:4]
```

They are explaining 88% of the total variance.

Factor Analysis

```
library(psych)
```

```
## Warning: package 'psych' was built under R version 3.5.3
```

```
#install.packages("psych", lib="/Library/Frameworks/R.framework/Versions/3.5/  
Resources/Library")
```

```
fit.pc <- principal(x_new, nfactors=4, rotate="varimax")
```

```
fit.pc
```

```
## Principal Components Analysis
```

```
## Call: principal(r = x_new, nfactors = 4, rotate = "varimax")
```

```
## Standardized loadings (pattern matrix) based upon correlation matrix
```

```
##          RC1   RC2   RC4   RC3   h2    u2 com
```

```
## clouds   0.90  0.12  0.10  0.01  0.83  0.17382 1.1
```

```
## pressure 0.01  0.95 -0.08  0.04  0.90  0.09521 1.0
```

```
## humidity 0.30  0.06  0.91  0.00  0.92  0.07603 1.2
```

```
## wind     0.13 -0.76 -0.50  0.00  0.85  0.15197 1.8
```

```
## distance 0.01  0.03  0.00  1.00  1.00  0.00015 1.0
```

```
## temp     0.81 -0.23  0.19  0.00  0.75  0.25060 1.3
```

```
##
```

```
##          RC1   RC2   RC4   RC3
```



```

## SS loadings          1.56 1.55 1.14 1.00
## Proportion Var      0.26 0.26 0.19 0.17
## Cumulative Var      0.26 0.52 0.71 0.88
## Proportion Explained 0.30 0.30 0.22 0.19
## Cumulative Proportion 0.30 0.59 0.81 1.00
##
## Mean item complexity = 1.2
## Test of the hypothesis that 4 components are sufficient.
##
## The root mean square of the residuals (RMSR) is 0.07
## with the empirical chi square 1476.35 with prob < NA
##
## Fit based upon off diagonal values = 0.93

round(fit.pc$values, 3)

## [1] 1.868 1.691 0.998 0.695 0.445 0.303

fit.pc$loadings

##
## Loadings:
##          RC1    RC2    RC4    RC3
## clouds    0.895  0.123
## pressure      0.947
## humidity  0.297      0.912
## wind      0.126 -0.761 -0.503
## distance      0.999
## temp      0.811 -0.233  0.195
##
##          RC1    RC2    RC4    RC3
## SS loadings  1.563 1.550 1.139 1.001
## Proportion Var 0.260 0.258 0.190 0.167
## Cumulative Var 0.260 0.519 0.709 0.875

```

The first 4 factors have an Eigenvalue >1 and which explains almost 88% of the variance. We can effectively reduce dimensionality from 6 to 4 while only losing about 11% of the variance.

Communalities

```

fit.pc$communality

## clouds pressure humidity wind distance temp
## 0.8261773 0.9047895 0.9239694 0.8480276 0.9998461 0.7493997

```

The variance in clouds accounted by all factors is 0.87, This is the extent to which an item correlates with all other items. All communalities are high, which means that the extracted components represent the variables well. If they are low, you may need to extract another component.


```
# Rotated factor scores, Notice the columns ordering: RC1, RC3, RC2 and RC4
fit.pc_scores<-fit.pc$scores
head(fit.pc_scores)
```

```
##           RC1      RC2      RC4      RC3
## [1,] 0.3196683 1.2040054 1.524236 0.05698478
## [2,] 0.3678396 0.9052307 1.329548 0.07812867
## [3,] 0.3823894 0.9322528 1.311205 -1.03514564
## [4,] 0.3872088 0.9242594 1.297225 -1.17642232
## [5,] 0.3902312 0.9315590 1.293890 -1.43719845
## [6,] 0.5484686 0.9711657 1.198470 -1.21705516
```

See factor recommendation

```
fa.parallel(x_new, fm='minres', fa='fa')
```

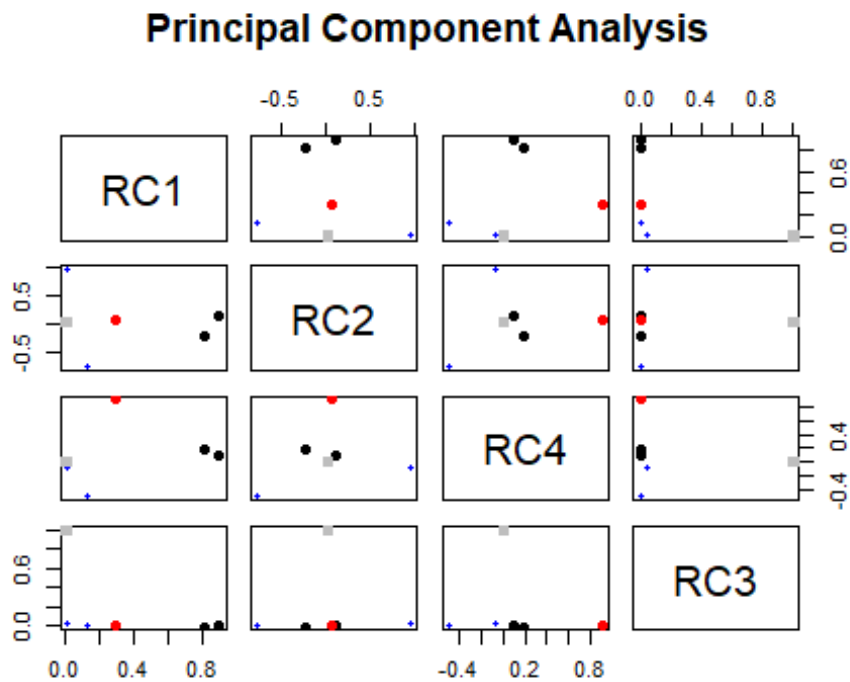
Parallel Analysis Scree Plots



```
## Parallel analysis suggests that the number of factors = 3 and the number
of components = NA
```

Blue line shows the eigen values of actual data and the two red lines show simulated and resampled data. Here factors between 2-4 will be a good choice..


```
fa.plot(fit.pc) # See Correlations within Factors
```



Visualize the relationship

```
fa.diagram(fit.pc)
```

Components Analysis



#Red dotted line

means Wind marginally falls under the RC1 bucket.


See Factor recommendations for a simple structure

```
vss(x_new)
```

```
## Warning in fac(r = r, nfactors = nfactors, n.obs = n.obs, rotate =  
## rotate, : A loading greater than abs(1) was detected. Examine the loadings  
## carefully.
```

```
## Warning in fa.stats(r = r, f = f, phi = phi, n.obs = n.obs, np.obs  
## = np.obs, : The estimated weights for the factor scores are probably  
## incorrect. Try a different factor extraction method.
```

```
## Warning in fac(r = r, nfactors = nfactors, n.obs = n.obs, rotate =  
## rotate, : An ultra-Heywood case was detected. Examine the results carefull  
y
```



```
##
## Very Simple Structure
## Call: vss(x = x_new)
## VSS complexity 1 achieves a maximum of 0.65 with 2 factors
## VSS complexity 2 achieves a maximum of 0.79 with 5 factors
##
## The Velicer MAP achieves a minimum of NA with 1 factors
## BIC achieves a minimum of NA with 2 factors
## Sample Size adjusted BIC achieves a minimum of NA with 2 factors
##
## Statistics by number of factors
##   vss1 vss2 map dof   chisq      prob sqresid fit RMSEA BIC SABIC
## 1 0.38 0.00 0.11  9 6.7e+03 0.0e+00    5.0 0.38 0.28 6597 6626
## 2 0.65 0.72 0.13  4 7.3e+02 3.3e-157    2.3 0.72 0.14  696  709
## 3 0.62 0.78 0.28  0 2.9e-01      NA    1.7 0.79  NA  NA  NA
## 4 0.55 0.74 0.44 -3 1.8e-06      NA    1.6 0.81  NA  NA  NA
## 5 0.64 0.79 1.00 -5 0.0e+00      NA    1.6 0.81  NA  NA  NA
## 6 0.64 0.79  NA -6 0.0e+00      NA    1.6 0.81  NA  NA  NA
##   complex eChisq SRMR eCRMS eBIC
## 1      1.0 9.6e+03 1.8e-01 0.239 9469
## 2      1.1 3.4e+02 3.5e-02 0.068  304
## 3      1.3 2.4e-01 9.2e-04  NA  NA
## 4      1.4 1.2e-06 2.1e-06  NA  NA
## 5      1.4 1.0e-17 6.1e-12  NA  NA
## 6      1.4 1.0e-17 6.1e-12  NA  NA
```

We should continue with 4 factors as it has the max vss complexity. #Regression analysis using the factors scores as the independent variable: Let's combine the dependent variable and the factor scores into a dataset and label them.

```
cab<-cbind(x[7],fit.pc$scores)
#Labelling the data
names(cab)<-c("Price","Wind_Pressure","Temperature","Humidity","Distance")
head(cab)
```

```
##      Price Wind_Pressure Temperature Humidity Distance
## 1 16.67376    0.3196683    1.2040054 1.524236  0.05698478
## 2 16.67376    0.3678396    0.9052307 1.329548  0.07812867
## 3  8.50000    0.3823894    0.9322528 1.311205 -1.03514564
## 4 16.50000    0.3872088    0.9242594 1.297225 -1.17642232
## 5 16.67376    0.3902312    0.9315590 1.293890 -1.43719845
## 6 26.50000    0.5484686    0.9711657 1.198470 -1.21705516
```

Let's split the dataset into training and testing dataset. (80:20)

```
set.seed(101)
Atrain<-sample(nrow(cab),nrow(cab)*0.80)
cab_train<-cab[Atrain,]
cab_test<-cab[-Atrain,]
dim(cab_train)

## [1] 7444    5

dim(cab_test)

## [1] 1862    5
```

Performing multiple regression (Taking alpha=0.1)

```
fit1 <- lm(Price~Wind_Pressure+Temperature+Humidity+Distance, data=cab_train)
#show the results
summary(fit1)
```

```
##
## Call:
## lm(formula = Price ~ Wind_Pressure + Temperature + Humidity +
##      Distance, data = cab_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -16.742   -0.392   -0.254   -0.063    69.318
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  16.683735   0.065216  255.823   <2e-16 ***
## Wind_Pressure  0.106653   0.065077   1.639    0.101
```

```
## Temperature    0.008777    0.065129    0.135    0.893
## Humidity       -0.036831    0.065151   -0.565    0.572
## Distance       1.919486    0.064994   29.533   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.627 on 7439 degrees of freedom
## Multiple R-squared:  0.1053, Adjusted R-squared:  0.1048
## F-statistic: 218.9 on 4 and 7439 DF,  p-value: < 2.2e-16
```

Wind_Pressure is highly insignificant

```
fit2 <- lm(Price~Temperature+Humidity+Distance, data=cab_train)
#show the results
summary(fit2)

##
## Call:
## lm(formula = Price ~ Temperature + Humidity + Distance, data = cab_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -16.768  -0.316  -0.244  -0.121   69.311
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  16.68374    0.06522  255.795   <2e-16 ***
## Temperature    0.00987    0.06513    0.152    0.880
## Humidity      -0.03725    0.06516   -0.572    0.568
## Distance       1.91947    0.06500   29.530   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.627 on 7440 degrees of freedom
## Multiple R-squared:  0.105, Adjusted R-squared:  0.1046
## F-statistic: 290.8 on 3 and 7440 DF,  p-value: < 2.2e-16
```

Humidity is insignificant

```
fit3 <- lm(Price~Temperature+Distance, data=cab_train)
#show the results
summary(fit3)


##
## Call:
## lm(formula = Price ~ Temperature + Distance, data = cab_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
```

```
## -16.779 -0.292 -0.233 -0.138 69.313
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 16.683592  0.065220 255.806  <2e-16 ***
## Temperature  0.009869  0.065130   0.152    0.88
## Distance     1.919638  0.064997  29.534  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.627 on 7441 degrees of freedom
## Multiple R-squared:  0.1049, Adjusted R-squared:  0.1047
## F-statistic: 436.1 on 2 and 7441 DF, p-value: < 2.2e-16
```

Now we can see that all the variables are significant now. This model is explaining only 0.1% of variance only.

Predicted Test

```
predicted_test <- predict(fit3, newdata= cab_test)
#Residuals Analysis
cab_testresults <- cab_test
cab_testresults$predicted <- predicted_test
cab_testresults$residual <- cab_testresults$Price - cab_testresults$predicted
#View(cab_testresults$residual)
plot(cab_testresults$predicted, cab_testresults$residual, xlab="Predicted", ylab="Residuals", pch=21, bg="red", col="red", main="Predicted Vs Residuals")
abline(0,0)
```

Anova Table

`anova(fit3)`

```
## Analysis of Variance Table
##
## Response: Price
##           Df Sum Sq Mean Sq  F value Pr(>F)
## Temperature    1      0      0.5    0.0143 0.9049
## Distance       1 27619 27619.1  872.2639 <2e-16 ***
## Residuals    7441 235610    31.7
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

`vcov(fit3)`

```
##           (Intercept)  Temperature      Distance
## (Intercept)  4.253612e-03 -9.381022e-06 -1.930894e-06
## Temperature -9.381022e-06  4.241977e-03  4.587644e-06
## Distance    -1.930894e-06  4.587644e-06  4.224652e-03
```

`cov2cor(vcov(fit3))`

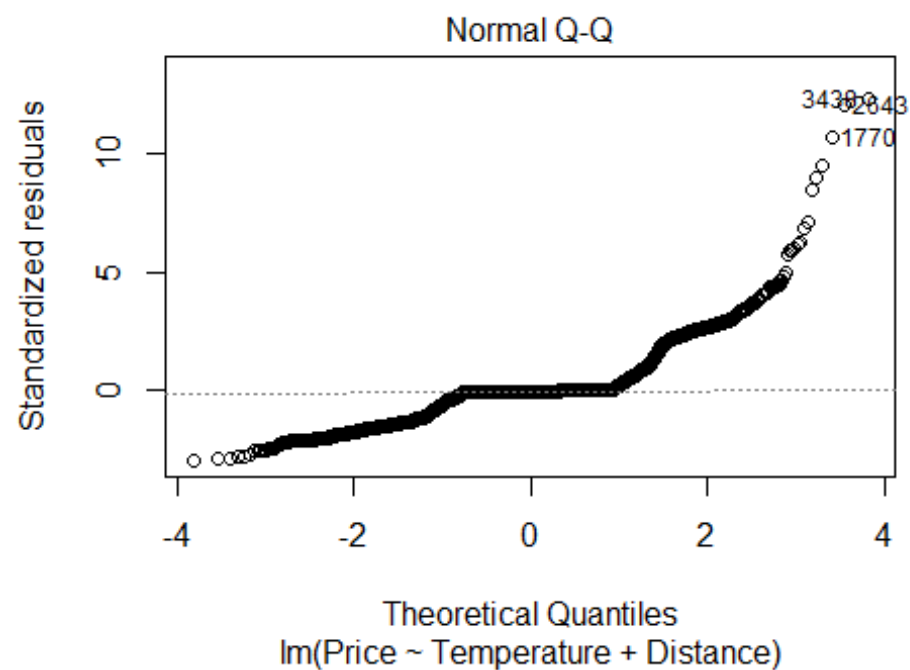
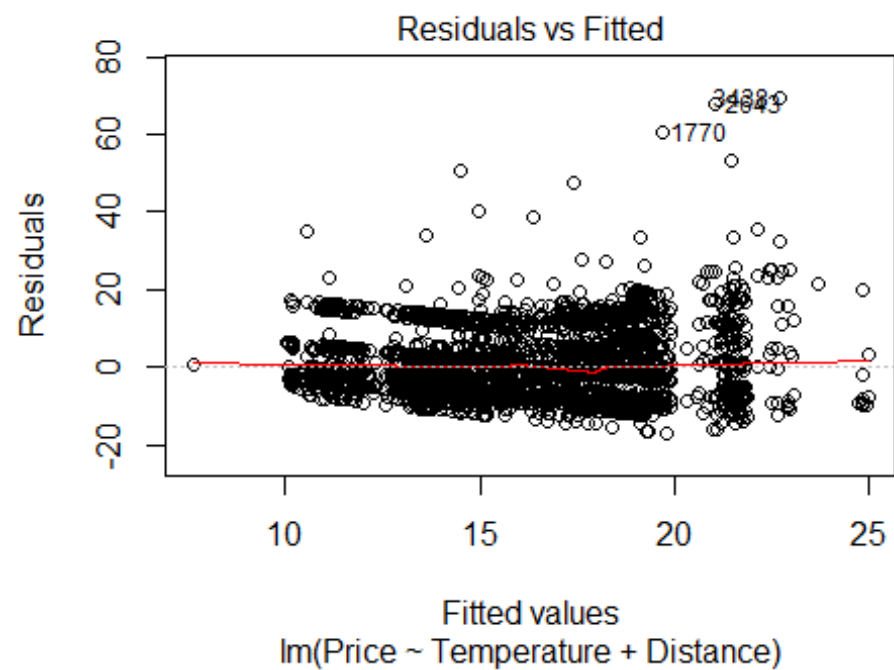
```
##           (Intercept)  Temperature      Distance
## (Intercept)  1.0000000000 -0.002208448 -0.0004554954
## Temperature -0.0022084476  1.0000000000  0.0010837025
## Distance    -0.0004554954  0.001083703  1.0000000000
```


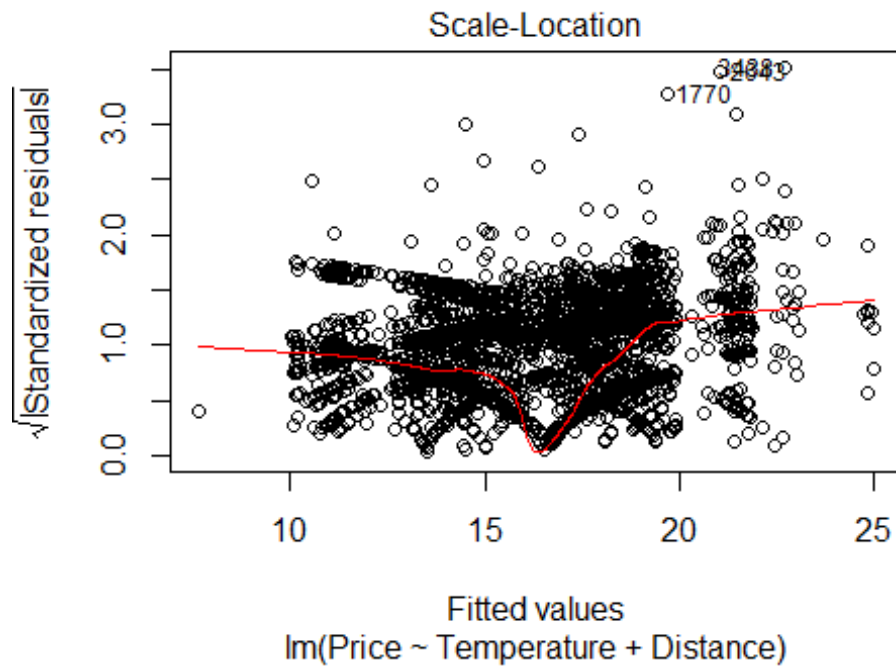
```
temp <- influence.measures(fit3)
head(temp)

## Influence measures of
##   lm(formula = Price ~ Temperature + Distance, data = cab_train) :
##
##           dfb.1_  dfb.Tmpr  dfb.Dstn      dffit cov.r   cook.d      hat inf
## 3464 -5.76e-04  4.35e-04 -8.20e-05 -7.26e-04 1.001 1.76e-07 0.000214
## 408  -6.28e-03  4.60e-03 -9.91e-03 -1.26e-02 1.001 5.30e-05 0.000542
## 6603 -4.95e-04  3.34e-04 -6.01e-05 -6.00e-04 1.001 1.20e-07 0.000198
## 6119 -3.17e-04 -4.89e-04 -2.17e-05 -5.84e-04 1.001 1.14e-07 0.000453
## 2325  2.44e-02  3.13e-02  1.83e-02  4.37e-02 0.999 6.36e-04 0.000429
## 2791 -3.83e-04 -3.84e-04 -3.34e-05 -5.44e-04 1.001 9.86e-08 0.000270
## 5440 -2.48e-04 -3.22e-04 -1.29e-05 -4.07e-04 1.001 5.53e-08 0.000360
```

diagnostic plots

```
plot(fit3)
```






Residuals Vs Fitted: Residuals are equally spread out around a horizontal line without distinct pattern that indicates we don't have non-linear relationships.

Normal Q-Q: The residuals are normally distributed if it has a straight linear line.

Scale-Location: Test homoscedasticity (assumption of equal variance) if you see a horizontal line with randomly spread points.

QQ-Plot

```
# Normality of Residuals
# qq plot for studentized resid
#install.packages("car")
qqPlot(fit3, main="QQ Plot")
```



```
## 2643 3438
## 643 2287
```


Distribution of studentized residuals

```
library(MASS)

## Warning: package 'MASS' was built under R version 3.5.3

sresid <- studres(fit3)
hist(sresid, freq=FALSE,
main="Distribution of Studentized Residuals")
xfit<-seq(min(sresid),max(sresid),length=40)
yfit<-dnorm(xfit)
lines(xfit, yfit)
```

Distribution of Studentized Residuals




#Hence we see that residuals are normally distributed

Cook's D plot

identify D values $> 4/(n-k-1)$

```
cutoff <- 4/((nrow(mtcars)-length(fit3$coefficients)-2))  
plot(fit3, which=4, cook.levels=cutoff)
```



Influence Plot

```
influencePlot(fit3, id.method="identify", main="Influence Plot", sub="Circle
size is proportional to Cook's Distance" )
```

```
## Warning in plot.window(...): "id.method" is not a graphical parameter
```

```
## Warning in plot.xy(xy, type, ...): "id.method" is not a graphical paramete
r
```

```
## Warning in axis(side = side, at = at, labels = labels, ...): "id.method" i
s
```

```
## not a graphical parameter
```

```
## Warning in axis(side = side, at = at, labels = labels, ...): "id.method" i
s
```


```
## not a graphical parameter
```

```
## Warning in box(...): "id.method" is not a graphical parameter
```

```
## Warning in title(...): "id.method" is not a graphical parameter
```

```
## Warning in plot.xy(xy.coords(x, y), type = type, ...): "id.method" is not
a
```

```
## graphical parameter
```



##	StudRes	Hat	CookD
## 2643	12.2042724	0.0008927121	4.349624e-02
## 3265	-1.3307620	0.0027551682	1.630725e-03
## 3438	12.4542355	0.0015229712	7.726163e-02
## 7475	0.1537065	0.0031021296	2.450922e-05


non-constant error variance test

```
ncvTest(fit3)
```

```
## Non-constant Variance Score Test
## Variance formula: ~ fitted.values
## Chisquare = 609.4441, Df = 1, p = < 2.22e-16
```

plot studentized residuals vs. fitted values

```
spreadLevelPlot(fit3)
```

```
##  
## Suggested power transformation: 3.193354
```